

import urllib

```
In [1]: import urllib
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [2]: housing = pd.read_csv('C:/articles/datasets/ml-master/ml-master/machine_learning/datasets/housing/housing.csv')
```

```
In [3]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   longitude              20640 non-null  float64
1   latitude              20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms        20433 non-null  float64
5   population             20640 non-null  float64
6   households             20640 non-null  float64
7   median_income          20640 non-null  float64
8   median_house_value     20640 non-null  float64
9   ocean_proximity        20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [4]: housing.describe()
```

Out[4]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_h...
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500

```
In [5]: housing['total_bedrooms'].value_counts()
#Thus the below output shows the distribution of the total bedrooms in the different blocks of california
```

Out[5]:

280.0	55
331.0	51
345.0	50
393.0	49
343.0	49
..	
2111.0	1
1852.0	1
1663.0	1
1652.0	1
2479.0	1

Name: total_bedrooms, Length: 1923, dtype: int64

```
In [6]: a = housing['housing_median_age'].value_counts()
#Thus the below output shows the distribution of the total bedrooms in the different blocks of california
```

```
In [7]: type(a)
```

Out[7]: pandas.core.series.Series

In [8]:


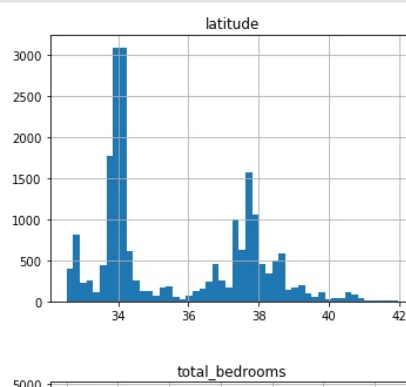
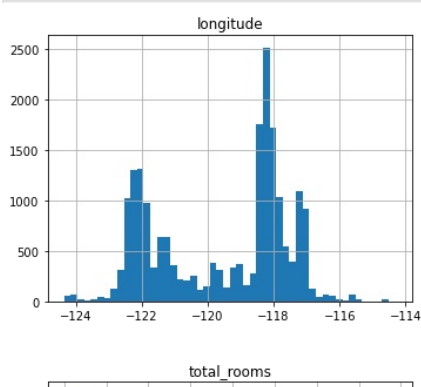
a







Out[8]:

```
52.0    1273
36.0     862
35.0     824
16.0     771
17.0     698
34.0     689
26.0     619
33.0     615
18.0     570
25.0     566
32.0     565
37.0     537
15.0     512
19.0     502
27.0     488
24.0     478
30.0     476
28.0     471
20.0     465
29.0     461
31.0     458
23.0     448
21.0     446
14.0     412
22.0     399
38.0     394
39.0     369
42.0     368
44.0     356
43.0     353
40.0     304
13.0     302
41.0     296
45.0     294
10.0     264
11.0     254
46.0     245
5.0      244
12.0     238
8.0      206
9.0      205
47.0     198
4.0      191
48.0     177
7.0      175
6.0      160
50.0     136
49.0     134
3.0       62
2.0       58
51.0      48
1.0        4
Name: housing_median_age, dtype: int64
```

In [9]:

```
#plotting the cloumn of each data frame..
%matplotlib inline
import matplotlib.pyplot as plt
housing.hist(bins = 50, figsize = (20,15))
plt.show()
```





```
In [10]: housing.head(5)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	N
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	N
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	N
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	N
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	N

```
In [11]: housing.describe()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_hu
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500

```
In [12]: housing['median_income'].value_counts()
```

```
Out[12]: 3.1250      49
15.0001      49
2.8750       46
2.6250       44
4.1250       44
..
5.0476        1
4.6379        1
2.9402        1
6.0256        1
5.5647        1
Name: median_income, Length: 12928, dtype: int64
```

```
In [13]: import random
random.random()
```

Out[13]: 0.25964032328767395

```
In [14]: b = np.array([1,2,7,4,6])
c =b[:3]
```

In [15]: c

Out[15]: array([1, 2, 7])

```
In [16]: np.random.rand(2,3)
```

Out[16]: array([[0.86431321, 0.23921777, 0.24903931],
[0.8429882 , 0.96657521, 0.87542146]])

```
In [17]: np.random.permutation(5)
```

Out[17]: array([4, 2, 3, 0, 1])

```
In [18]: np.random.permutation(5)
```

Out[18]: array([3, 0, 1, 4, 2])

```
In [19]: #splitting the dataset into training and test set
np.random.seed(42)
def split_train_test(data, test_ratio):
    shuffled_indices = np.random.permutation(len(data))
    test_set_size =int(len(data)*test_ratio)
    test_indices = shuffled_indices[:test_set_size]
    train_indices = shuffled_indices[test_set_size:]
    return data.iloc[train_indices] ,data.iloc[test_indices]
train_set , test_set = split_train_test(housing, 0.2)
print(len(train_set) , "train +" , len(test_set) , "test")
```

16512 train + 4128 test

```
In [20]: train_set.head()
```

Out[20]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
14196	-117.03	32.71	33.0	3126.0	627.0	2300.0	623.0	3.2596	103000.0	
8267	-118.16	33.77	49.0	3382.0	787.0	1314.0	756.0	3.8125	382100.0	
17445	-120.48	34.66	4.0	1897.0	331.0	915.0	336.0	4.1563	172600.0	
14265	-117.11	32.69	36.0	1421.0	367.0	1418.0	355.0	1.9425	93400.0	
2271	-119.80	36.78	43.0	2382.0	431.0	874.0	380.0	3.5542	96500.0	

```
In [21]: test_set.head()
```

Out[21]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
20046	-119.01	36.06	25.0	1505.0	NaN	1392.0	359.0	1.6812	47700.0	
3024	-119.46	35.14	30.0	2943.0	NaN	1565.0	584.0	2.5313	45800.0	
15663	-122.44	37.80	52.0	3830.0	NaN	1310.0	963.0	3.4801	500001.0	
20484	-118.72	34.28	17.0	3051.0	NaN	1705.0	495.0	5.7376	218600.0	
9814	-121.93	36.62	34.0	2351.0	NaN	1063.0	428.0	3.7250	278000.0	

```
In [22]: from sklearn.model_selection import train_test_split
```

```
In [23]: train_set , test_set = train_test_split(housing, test_size =0.2, random_state =42)
```

```
In [24]: print(len(train_set), 'train +' ,len(test_set), 'test')
```

```
16512 train + 4128 test
```


```
In [25]: test_set.head()
```

```
Out[25]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
20046	-119.01	36.06	25.0	1505.0	NaN	1392.0	359.0	1.6812	47700.0	
3024	-119.46	35.14	30.0	2943.0	NaN	1565.0	584.0	2.5313	45800.0	
15663	-122.44	37.80	52.0	3830.0	NaN	1310.0	963.0	3.4801	500001.0	
20484	-118.72	34.28	17.0	3051.0	NaN	1705.0	495.0	5.7376	218600.0	
9814	-121.93	36.62	34.0	2351.0	NaN	1063.0	428.0	3.7250	278000.0	

```
In [26]: housing['median_income'].hist()
```

```
Out[26]: <AxesSubplot:>
```



**from the graph of the median income, it is evident that most of the population lies in the range 2 to 5. Thus we need to make our strata from this portion of the distribution.

```
In [27]: housing['income_cat'] = np.ceil(housing['median_income']/1.5)
```


```
In [28]: housing['income_cat'].value_counts()
```

```
Out[28]: 3.0    7236
2.0    6581
4.0    3639
5.0    1423
1.0     822
6.0     532
7.0     189
8.0     105
9.0      50
11.0     49
10.0     14
Name: income_cat, dtype: int64
```

```
In [29]: housing['income_cat'].hist()
```

```
Out[29]: <AxesSubplot:>
```





```
In [30]: housing['median_income'].value_counts()
```

```
Out[30]: 3.1250      49
15.0001      49
2.8750       46
2.6250       44
4.1250       44
..
5.0476        1
4.6379        1
2.9402        1
6.0256        1
5.5647        1
Name: median_income, Length: 12928, dtype: int64
```


```
In [31]: housing['income_cat'].where(housing['income_cat']<5 , 5.0, inplace = True)
```

```
In [32]: housing['income_cat'].value_counts()
```

```
Out[32]: 3.0      7236
2.0      6581
4.0      3639
5.0      2362
1.0       822
Name: income_cat, dtype: int64
```

```
In [33]: housing['income_cat'].hist()
```

```
Out[33]: <AxesSubplot:>
```



```
In [34]: housing['income_cat'].describe()
```

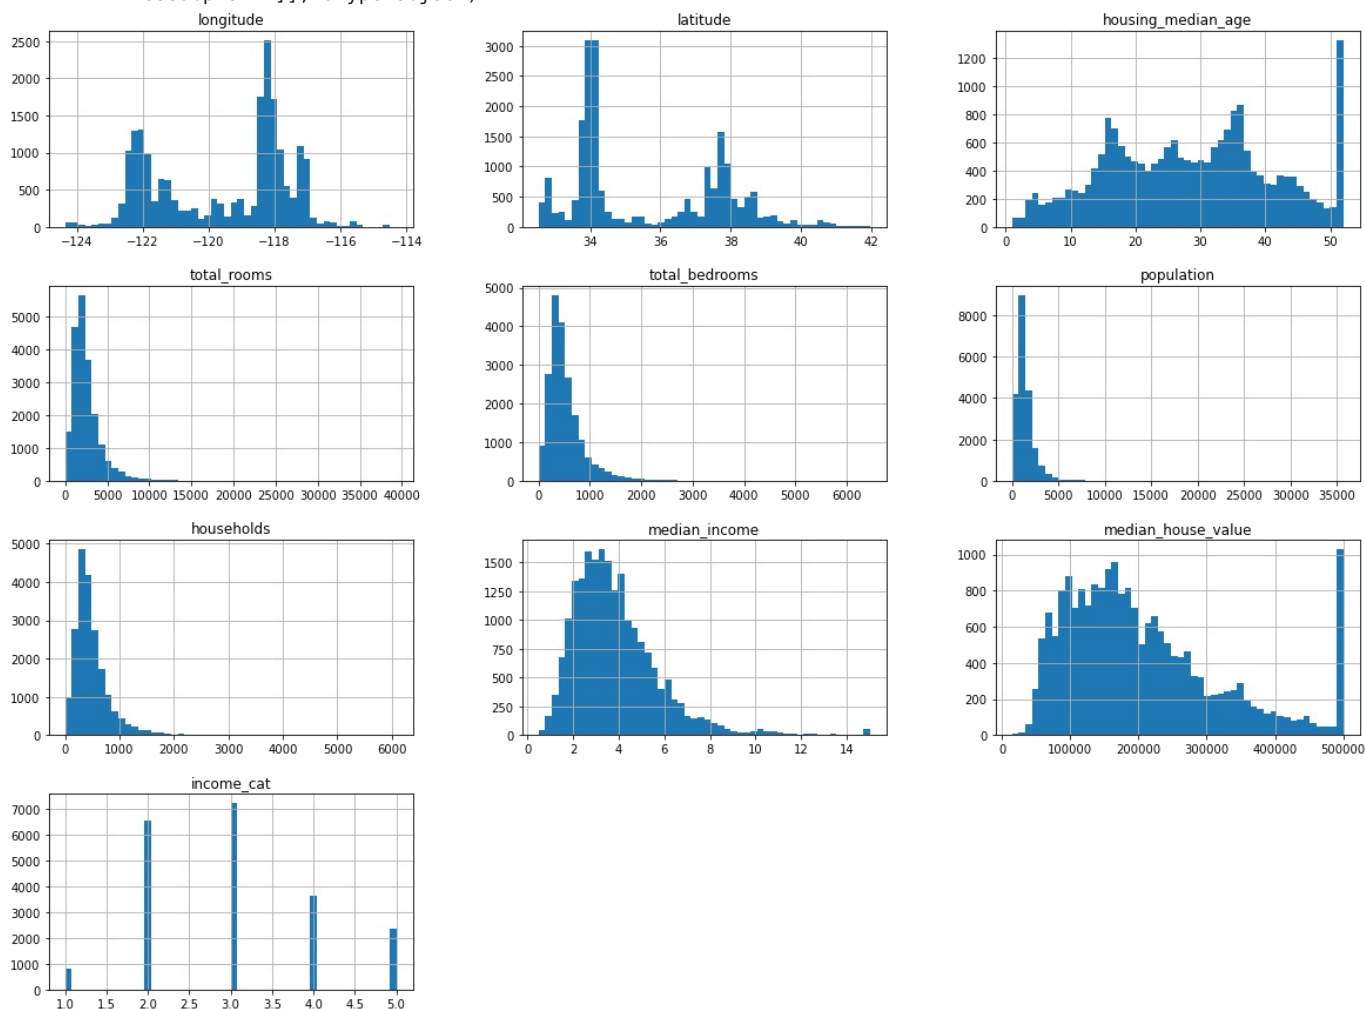
```
Out[34]: count      20640.000000
mean          3.006686
std           1.054618
min           1.000000
25%           2.000000
50%           3.000000
75%           4.000000
max           5.000000
Name: income_cat, dtype: float64
```

```
In [35]: housing['median_income'].describe()
```

```
Out[35]: count    20640.000000
mean         3.870671
std          1.899822
min          0.499900
25%          2.563400
50%          3.534800
75%          4.743250
max          15.000100
Name: median_income, dtype: float64
```


```
In [36]: housing.hist(bins =50 ,figsize =(20,15))
```


```
Out[36]: array([[<AxesSubplot:title={ 'center': 'longitude'>,
  <AxesSubplot:title={ 'center': 'latitude'>,
  <AxesSubplot:title={ 'center': 'housing_median_age'>],
  [<AxesSubplot:title={ 'center': 'total_rooms'>,
  <AxesSubplot:title={ 'center': 'total_bedrooms'>,
  <AxesSubplot:title={ 'center': 'population'>],
  [<AxesSubplot:title={ 'center': 'households'>,
  <AxesSubplot:title={ 'center': 'median_income'>,
  <AxesSubplot:title={ 'center': 'median_house_value'>],
  [<AxesSubplot:title={ 'center': 'income_cat'>, <AxesSubplot:>,
  <AxesSubplot:>]], dtype=object)
```



```
In [37]: housing['median_income'].hist()
```

```
Out[37]: <AxesSubplot:>
```





```
In [38]: housing['median_income'].value_counts()
```

```
Out[38]: 3.1250    49
         15.0001   49
         2.8750   46
         2.6250   44
         4.1250   44
         ..
         5.0476    1
         4.6379    1
         2.9402    1
         6.0256    1
         5.5647    1
Name: median_income, Length: 12928, dtype: int64
```

```
In [39]: housing['income_cat'].value_counts()
```

```
Out[39]: 3.0    7236
         2.0    6581
         4.0    3639
         5.0    2362
         1.0     822
Name: income_cat, dtype: int64
```

```
In [40]: len(housing)
```

```
Out[40]: 20640
```

```
In [41]: from sklearn.model_selection import StratifiedShuffleSplit
```

```
In [42]: split= StratifiedShuffleSplit(n_splits=1, test_size =0.2,random_state=42)
```

```
In [43]: split
```

```
Out[43]: StratifiedShuffleSplit(n_splits=1, random_state=42, test_size=0.2,
                                train_size=None)
```

```
In [44]: for train_index , test_index in split.split(housing, housing['income_cat']):
         strat_train_set = housing.loc[train_index]
         strat_test_set = housing.loc[test_index]
```

```
In [45]: strat_train_set.head()
```

```
Out[45]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	286600.0	
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	340600.0	
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	196900.0	
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	46300.0	
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	254500.0	


```
In [46]: strat_train_set['income_cat'].value_counts()
```

```
Out[46]: 3.0    5789
         2.0    5265
         4.0    2911
         5.0    1889
         1.0     658
         Name: income_cat, dtype: int64
```

```
In [47]: strat_test_set['income_cat'].value_counts()
```

```
Out[47]: 3.0    1447
         2.0    1316
         4.0     728
         5.0     473
         1.0     164
         Name: income_cat, dtype: int64
```

```
In [48]: strat_test_set.head()
```

```
Out[48]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
5241	-118.39	34.12	29.0	6447.0	1012.0	2184.0	960.0	8.2816	500001.0	
10970	-117.86	33.77	39.0	4159.0	655.0	1669.0	651.0	4.6111	240300.0	
20351	-119.05	34.21	27.0	4357.0	926.0	2110.0	876.0	3.0119	218200.0	
6568	-118.15	34.20	52.0	1786.0	306.0	1018.0	322.0	4.1518	182100.0	
13285	-117.68	34.07	32.0	1775.0	314.0	1067.0	302.0	4.0375	121300.0	

```
In [49]: strat_test_set.head()
```

```
Out[49]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
5241	-118.39	34.12	29.0	6447.0	1012.0	2184.0	960.0	8.2816	500001.0	
10970	-117.86	33.77	39.0	4159.0	655.0	1669.0	651.0	4.6111	240300.0	
20351	-119.05	34.21	27.0	4357.0	926.0	2110.0	876.0	3.0119	218200.0	
6568	-118.15	34.20	52.0	1786.0	306.0	1018.0	322.0	4.1518	182100.0	
13285	-117.68	34.07	32.0	1775.0	314.0	1067.0	302.0	4.0375	121300.0	

```
In [50]: # Income category proportion in test set generated with stratified sampling
strat_test_set["income_cat"].value_counts() / len(strat_test_set)
```

```
Out[50]: 3.0    0.350533
         2.0    0.318798
         4.0    0.176357
         5.0    0.114583
         1.0    0.039729
         Name: income_cat, dtype: float64
```

****This shows that 35% of the people have income upto 3000 dollars and 31% of them have income upto 2000 dollars**

```
In [51]: #this piece of code removes the extra column we created before
for set_in (strat_train_set , strat_test_set):
    set_.drop('income_cat' , axis =1 , inplace=True)
```

```
In [52]: strat_test_set.head()
```

```
Out[52]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
5241	-118.39	34.12	29.0	6447.0	1012.0	2184.0	960.0	8.2816	500001.0	

10970	-117.86	33.77	39.0	4159.0	655.0	1669.0	651.0	4.6111	240300.0
20351	-119.05	34.21	27.0	4357.0	926.0	2110.0	876.0	3.0119	218200.0
6568	-118.15	34.20	52.0	1786.0	306.0	1018.0	322.0	4.1518	182100.0
13285	-117.68	34.07	32.0	1775.0	314.0	1067.0	302.0	4.0375	121300.0

```
In [53]: strat_train_set.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	286600.0	
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	340600.0	
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	196900.0	
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	46300.0	
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	254500.0	

```
In [54]: traindata = strat_train_set.copy()
```


```
In [55]: traindata
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	286600.0	
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	340600.0	
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	196900.0	
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	46300.0	
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	254500.0	
...	
6563	-118.13	34.20	46.0	1271.0	236.0	573.0	210.0	4.9312	240200.0	
12053	-117.56	33.88	40.0	1196.0	294.0	1052.0	258.0	2.0682	113000.0	
13908	-116.40	34.09	9.0	4855.0	872.0	2098.0	765.0	3.2723	97800.0	
11159	-118.01	33.82	31.0	1960.0	380.0	1356.0	356.0	4.0625	225900.0	
15775	-122.45	37.77	52.0	3095.0	682.0	1269.0	639.0	3.5750	500001.0	

16512 rows × 10 columns


```
In [56]: #visualisation of the data
traindata.plot(kind='scatter', x='longitude', y='latitude')
```

```
Out[56]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```




```
In [57]: #visualisation of the data
traindata.plot(kind='scatter', x='longitude', y='latitude', alpha =0.1)
```

```
Out[57]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```




```
In [58]: #visualisation of the data
traindata.plot(kind='scatter', x='longitude', y='latitude', alpha=.4)
```

```
Out[58]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```




```
In [59]: traindata.plot(kind='scatter', x='longitude', y='latitude', alpha=0.5, s=traindata['population']/100, label='population',
plt.legend())
```

```
Out[59]: <matplotlib.legend.Legend at 0x1a68aface20>
```



```
In [60]: traindata.plot(kind='scatter', x='longitude', y='latitude', s=traindata['population']/100, label='population',
plt.legend())
```

```
Out[60]: <matplotlib.legend.Legend at 0x1a68ae92ca0>
```



```
In [61]: #finding correlation between the different variables at play
corr_matrix = traindata.corr()
corr_matrix
```

```
Out[61]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_
longitude	1.000000	-0.924478	-0.105848	0.048871	0.076598	0.108030	0.063070	-0.019583	
latitude	-0.924478	1.000000	0.005766	-0.039184	-0.072419	-0.115222	-0.077647	-0.075205	
housing_median_age	-0.105848	0.005766	1.000000	-0.364509	-0.325047	-0.298710	-0.306428	-0.111360	
total_rooms	0.048871	-0.039184	-0.364509	1.000000	0.929379	0.855109	0.918392	0.200087	
total_bedrooms	0.076598	-0.072419	-0.325047	0.929379	1.000000	0.876320	0.980170	-0.009740	
population	0.108030	-0.115222	-0.298710	0.855109	0.876320	1.000000	0.904637	0.002380	
households	0.063070	-0.077647	-0.306428	0.918392	0.980170	0.904637	1.000000	0.010781	
median_income	-0.019583	-0.075205	-0.111360	0.200087	-0.009740	0.002380	0.010781	1.000000	
median_house_value	-0.047432	-0.142724	0.114110	0.135097	0.047689	-0.026920	0.064506	0.687160	

```
In [62]: corr_matrix['median_house_value']
```

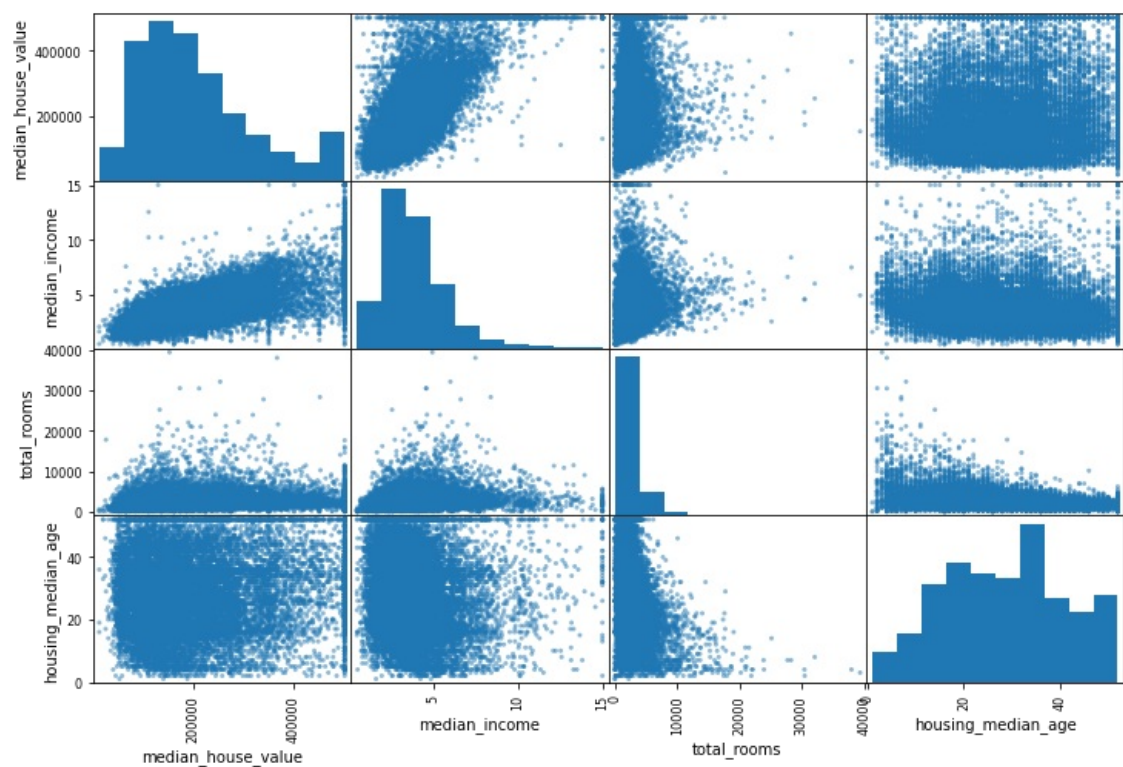
```
Out[62]: longitude      -0.047432
latitude      -0.142724
housing_median_age  0.114110
total_rooms     0.135097
total_bedrooms  0.047689
population     -0.026920
households      0.064506
median_income   0.687160
median_house_value 1.000000
Name: median_house_value, dtype: float64
```

```
In [63]: #to sort the corr_matrix in descending order
corr_matrix['median_house_value'].sort_values(ascending =False)
```

```
Out[63]: median_house_value  1.000000
median_income   0.687160
total_rooms     0.135097
housing_median_age  0.114110
households      0.064506
total_bedrooms  0.047689
population     -0.026920
longitude      -0.047432
latitude      -0.142724
Name: median_house_value, dtype: float64
```

```
In [64]: #visualisation
from pandas.plotting import scatter_matrix
attributes = ['median_house_value', 'median_income', 'total_rooms', 'housing_median_age']
scatter_matrix(traindata[attributes], figsize=(12,8))
```

```
Out[64]: array([[<AxesSubplot:xlabel='median_house_value', ylabel='median_house_value'>,
<AxesSubplot:xlabel='median_income', ylabel='median_house_value'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_house_value'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_house_value'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='median_income'>,
<AxesSubplot:xlabel='median_income', ylabel='median_income'>,
<AxesSubplot:xlabel='total_rooms', ylabel='median_income'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='median_income'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='total_rooms'>,
<AxesSubplot:xlabel='median_income', ylabel='total_rooms'>,
<AxesSubplot:xlabel='total_rooms', ylabel='total_rooms'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='total_rooms'>],
[<AxesSubplot:xlabel='median_house_value', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='median_income', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='total_rooms', ylabel='housing_median_age'>,
<AxesSubplot:xlabel='housing_median_age', ylabel='housing_median_age'>]],
dtype=object)
```



```
In [65]: traindata[['median_income', 'median_house_value']].head()
```

Out[65]:

	median_income	median_house_value
17606	2.7042	286600.0
18632	6.4214	340600.0
14650	2.8621	196900.0
3230	1.8839	46300.0
3555	3.0347	254500.0

```
In [66]: traindata.describe()
```

Out[66]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_h
count	16512.000000	16512.000000	16512.000000	16512.000000	16354.000000	16512.000000	16512.000000	16512.000000	16
mean	-119.575834	35.639577	28.653101	2622.728319	534.973890	1419.790819	497.060380	3.875589	206
std	2.001860	2.138058	12.574726	2138.458419	412.699041	1115.686241	375.720845	1.904950	115
min	-124.350000	32.540000	1.000000	6.000000	2.000000	3.000000	2.000000	0.499900	14

25%	-121.800000	33.940000	18.000000	1443.000000	295.000000	784.000000	279.000000	2.566775	119
50%	-118.510000	34.260000	29.000000	2119.500000	433.000000	1164.000000	408.000000	3.540900	179
75%	-118.010000	37.720000	37.000000	3141.000000	644.000000	1719.250000	602.000000	4.744475	263
max	-114.310000	41.950000	52.000000	39320.000000	6210.000000	35682.000000	5358.000000	15.000100	500

```
In [67]: traindata['median_income'].value_counts()
```

```
Out[67]: 15.0001    38
          2.8750    38
          2.6250    37
          3.1250    37
          3.3750    33
          ..
          1.6410     1
          6.0074     1
          3.8421     1
          2.6941     1
          5.5647     1
Name: median_income, Length: 10905, dtype: int64
```


```
In [68]: traindata.describe()
```

```
Out[68]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_hu
count	16512.000000	16512.000000	16512.000000	16512.000000	16354.000000	16512.000000	16512.000000	16512.000000	16
mean	-119.575834	35.639577	28.653101	2622.728319	534.973890	1419.790819	497.060380	3.875589	206
std	2.001860	2.138058	12.574726	2138.458419	412.699041	1115.686241	375.720845	1.904950	115
min	-124.350000	32.540000	1.000000	6.000000	2.000000	3.000000	2.000000	0.499900	14
25%	-121.800000	33.940000	18.000000	1443.000000	295.000000	784.000000	279.000000	2.566775	119
50%	-118.510000	34.260000	29.000000	2119.500000	433.000000	1164.000000	408.000000	3.540900	179
75%	-118.010000	37.720000	37.000000	3141.000000	644.000000	1719.250000	602.000000	4.744475	263
max	-114.310000	41.950000	52.000000	39320.000000	6210.000000	35682.000000	5358.000000	15.000100	500


```
In [69]: traindata['median_house_value'].hist()
```


```
Out[69]: <AxesSubplot:>
```



```
In [70]: traindata.plot(kind='scatter', x='median_income', y='median_house_value')
```

```
Out[70]: <AxesSubplot:xlabel='median_income', ylabel='median_house_value'>
```





***The graph above shows the following observation:

- 1) There exist a positive correlation between the median income and the median house values.
- 2) The median house value above 50K has been capped at 50K.

```
In [71]: #creating new attributes
traindata['rooms_per_households']= traindata['total_rooms']/traindata['households']
traindata['bedrooms_per_room']= traindata['total_bedrooms']/traindata['total_rooms']
traindata['population_per_houselholds']= traindata['population']/traindata['households']
```

```
In [72]: traindata
```

```
Out[72]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	oc
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	286600.0	
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	340600.0	
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	196900.0	
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	46300.0	
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	254500.0	
...
6563	-118.13	34.20	46.0	1271.0	236.0	573.0	210.0	4.9312	240200.0	
12053	-117.56	33.88	40.0	1196.0	294.0	1052.0	258.0	2.0682	113000.0	
13908	-116.40	34.09	9.0	4855.0	872.0	2098.0	765.0	3.2723	97800.0	
11159	-118.01	33.82	31.0	1960.0	380.0	1356.0	356.0	4.0625	225900.0	
15775	-122.45	37.77	52.0	3095.0	682.0	1269.0	639.0	3.5750	500001.0	

16512 rows × 13 columns


```
In [73]: #finding correlation between the diffrent variables at play
new_corr_matrix = traindata.corr()
```


```
In [74]: #to sort the corr_matrix in descending order
new_corr_matrix['median_house_value'].sort_values(ascending =False)
```

```
Out[74]: median_house_value      1.000000
median_income      0.687160
rooms_per_households 0.146285
total_rooms      0.135097
housing_median_age  0.114110
households      0.064506
total_bedrooms     0.047689
population_per_houselholds -0.021985
population      -0.026920
longitude         -0.047432
latitude         -0.142724
bedrooms_per_room  -0.259984
Name: median_house_value, dtype: float64
```

```
In [75]: traindata.plot(kind='scatter', x='bedrooms_per_room', y='median_house_value', alpha=.4)
```

```
Out[75]: <AxesSubplot:xlabel='bedrooms_per_room', ylabel='median_house_value'>
```





```
In [76]: # Let's revert to a clean training set

traindata = strat_train_set.drop("median_house_value", axis=1) # drop labels for training set
traindata_labels = strat_train_set["median_house_value"].copy()

# Note drop() creates a copy of the data and does not affect strat_train_set
```

```
In [77]: traindata.head(5)
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN

```
In [78]: traindata_labels.head()
```

```
Out[78]: 17606    286600.0
18632    340600.0
14650    196900.0
3230     46300.0
3555    254500.0
Name: median_house_value, dtype: float64
```

```
In [79]: traindata.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 16512 entries, 17606 to 15775
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   longitude              16512 non-null  float64
1   latitude               16512 non-null  float64
2   housing_median_age     16512 non-null  float64
3   total_rooms            16512 non-null  float64
4   total_bedrooms         16354 non-null  float64
5   population             16512 non-null  float64
6   households              16512 non-null  float64
7   median_income          16512 non-null  float64
8   ocean_proximity        16512 non-null  object
dtypes: float64(8), object(1)
memory usage: 1.3+ MB
```

```
In [80]: housing.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   longitude              20640 non-null  float64
1   latitude               20640 non-null  float64
2   housing_median_age     20640 non-null  float64
3   total_rooms            20640 non-null  float64
4   total_bedrooms         20433 non-null  float64
```



```
5    population      20640 non-null float64
6    households      20640 non-null float64
7    median_income    20640 non-null float64
8    median_house_value 20640 non-null float64
9    ocean_proximity  20640 non-null object
10   income_cat       20640 non-null float64
dtypes: float64(10), object(1)
memory usage: 1.7+ MB
```

```
In [81]: traindata['total_bedrooms'].isna()
```

```
Out[81]: 17606    False
18632    False
14650    False
3230     False
3555     False
...
6563     False
12053    False
13908    False
11159    False
15775    False
Name: total_bedrooms, Length: 16512, dtype: bool
```

```
In [82]: isn = traindata.isna()
isn.any(axis=1)
```

```
Out[82]: 17606    False
18632    False
14650    False
3230     False
3555     False
...
6563     False
12053    False
13908    False
11159    False
15775    False
Length: 16512, dtype: bool
```

```
In [83]: type(isn)
```

```
Out[83]: pandas.core.frame.DataFrame
```

```
In [84]: sample_incomplete_rows = traindata[traindata.isna().any(axis =1)]
sample_incomplete_rows
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
4629	-118.30	34.07	18.0	3759.0	NaN	3296.0	1462.0	2.2708	<1H OCEAN
6068	-117.86	34.01	16.0	4632.0	NaN	3038.0	727.0	5.1762	<1H OCEAN
17923	-121.97	37.35	30.0	1955.0	NaN	999.0	386.0	4.6328	<1H OCEAN
13656	-117.30	34.05	6.0	2155.0	NaN	1039.0	391.0	1.6675	INLAND
19252	-122.79	38.48	7.0	6837.0	NaN	3468.0	1405.0	3.1662	<1H OCEAN
...
3376	-118.28	34.25	29.0	2559.0	NaN	1886.0	769.0	2.6036	<1H OCEAN
4691	-118.37	34.07	50.0	2519.0	NaN	1117.0	516.0	4.3667	<1H OCEAN
6052	-117.76	34.04	34.0	1914.0	NaN	1564.0	328.0	2.8347	INLAND
17198	-119.75	34.45	6.0	2864.0	NaN	1404.0	603.0	5.5073	NEAR OCEAN
4738	-118.38	34.05	49.0	702.0	NaN	458.0	187.0	4.8958	<1H OCEAN

158 rows × 9 columns

```
In [85]: sample_incomplete_rows.dropna(subset=['total_bedrooms'])
```

```
Out[85]: longitude latitude housing_median_age total_rooms total_bedrooms population households median_income ocean_proximity
```

```
In [86]: ## Or let's drop the column of the missing values
sample_incomplete_rows.drop('total_bedrooms',axis =1)
```

```
Out[86]:
```

	longitude	latitude	housing_median_age	total_rooms	population	households	median_income	ocean_proximity
4629	-118.30	34.07	18.0	3759.0	3296.0	1462.0	2.2708	<1H OCEAN
6068	-117.86	34.01	16.0	4632.0	3038.0	727.0	5.1762	<1H OCEAN
17923	-121.97	37.35	30.0	1955.0	999.0	386.0	4.6328	<1H OCEAN
13656	-117.30	34.05	6.0	2155.0	1039.0	391.0	1.6675	INLAND
19252	-122.79	38.48	7.0	6837.0	3468.0	1405.0	3.1662	<1H OCEAN
...
3376	-118.28	34.25	29.0	2559.0	1886.0	769.0	2.6036	<1H OCEAN
4691	-118.37	34.07	50.0	2519.0	1117.0	516.0	4.3667	<1H OCEAN
6052	-117.76	34.04	34.0	1914.0	1564.0	328.0	2.8347	INLAND
17198	-119.75	34.45	6.0	2864.0	1404.0	603.0	5.5073	NEAR OCEAN
4738	-118.38	34.05	49.0	702.0	458.0	187.0	4.8958	<1H OCEAN

158 rows × 8 columns

```
In [87]: median = traindata['total_bedrooms'].median()
median
```

```
Out[87]: 433.0
```

```
In [88]: sample_incomplete_rows['total_bedrooms'].fillna(median, inplace=True)
```

C:\Users\NIYAZ AHMED\anaconda3\lib\site-packages\pandas\core\series.py:4463: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
return super().fillna()

```
In [89]: sample_incomplete_rows
```

```
Out[89]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
4629	-118.30	34.07	18.0	3759.0	433.0	3296.0	1462.0	2.2708	<1H OCEAN
6068	-117.86	34.01	16.0	4632.0	433.0	3038.0	727.0	5.1762	<1H OCEAN
17923	-121.97	37.35	30.0	1955.0	433.0	999.0	386.0	4.6328	<1H OCEAN
13656	-117.30	34.05	6.0	2155.0	433.0	1039.0	391.0	1.6675	INLAND
19252	-122.79	38.48	7.0	6837.0	433.0	3468.0	1405.0	3.1662	<1H OCEAN
...
3376	-118.28	34.25	29.0	2559.0	433.0	1886.0	769.0	2.6036	<1H OCEAN
4691	-118.37	34.07	50.0	2519.0	433.0	1117.0	516.0	4.3667	<1H OCEAN
6052	-117.76	34.04	34.0	1914.0	433.0	1564.0	328.0	2.8347	INLAND
17198	-119.75	34.45	6.0	2864.0	433.0	1404.0	603.0	5.5073	NEAR OCEAN
4738	-118.38	34.05	49.0	702.0	433.0	458.0	187.0	4.8958	<1H OCEAN

158 rows × 9 columns

```
In [90]: traindata.head()
```

Out[90]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN

In [91]:

```
sample_incomplete_rows.iloc[1]
```

Out[91]:

```
longitude      -117.86
latitude        34.01
housing_median_age    16.0
total_rooms      4632.0
total_bedrooms    433.0
population       3038.0
households       727.0
median_income     5.1762
ocean_proximity    <1H OCEAN
Name: 6068, dtype: object
```

In [92]:

```
traindata.iloc[4629]
```

Out[92]:

```
longitude      -119.81
latitude        36.73
housing_median_age    47.0
total_rooms     1314.0
total_bedrooms    416.0
population      1155.0
households       326.0
median_income     1.372
ocean_proximity    INLAND
Name: 1987, dtype: object
```

In [93]:

```
# Let's use Scikit-Learn Imputer class to fill missing values

from sklearn.impute import SimpleImputer
imputer = SimpleImputer(strategy='median')
```

In [94]:

```
traindata_num = traindata.drop('ocean_proximity', axis =1)
```

In [95]:

```
traindata_num
```

Out[95]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347
...
6563	-118.13	34.20	46.0	1271.0	236.0	573.0	210.0	4.9312
12053	-117.56	33.88	40.0	1196.0	294.0	1052.0	258.0	2.0682
13908	-116.40	34.09	9.0	4855.0	872.0	2098.0	765.0	3.2723
11159	-118.01	33.82	31.0	1960.0	380.0	1356.0	356.0	4.0625
15775	-122.45	37.77	52.0	3095.0	682.0	1269.0	639.0	3.5750

16512 rows × 8 columns

In [96]:

```
traindata_num[traindata_num.isna().any(axis = 1)].head()
```

Out[96]:

```
longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households  median_income
```

4629	-118.30	34.07	18.0	3759.0	NaN	3296.0	1462.0	2.2708
6068	-117.86	34.01	16.0	4632.0	NaN	3038.0	727.0	5.1762
17923	-121.97	37.35	30.0	1955.0	NaN	999.0	386.0	4.6328
13656	-117.30	34.05	6.0	2155.0	NaN	1039.0	391.0	1.6675
19252	-122.79	38.48	7.0	6837.0	NaN	3468.0	1405.0	3.1662

```
In [97]: traindata_num.head()
```

```
Out[97]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347

```
In [98]: # creating an object of the SimpleImputer class
imputer = SimpleImputer(strategy = 'median')
```

```
In [99]: imputer.fit(traindata_num)
```

```
Out[99]: SimpleImputer(strategy='median')
```

```
In [100]: traindata.median()
```

```
Out[100]: longitude      -118.5100
latitude        34.2600
housing_median_age  29.0000
total_rooms      2119.5000
total_bedrooms    433.0000
population        1164.0000
households        408.0000
median_income      3.5409
dtype: float64
```

```
In [101]: imputer.statistics_
```

```
Out[101]: array([[-118.51 ,  34.26 ,  29.    , 2119.5 ,  433.    , 1164.    ,
                   408.    ,  3.5409])
```

```
In [102]: X = imputer.transform(traindata_num)
```

```
In [103]: X
```

```
Out[103]: array([[ -121.89 ,  37.29 ,  38.    , ...,  710.    ,  339.    ,
                    2.7042],
                  [ -121.93 ,  37.05 ,  14.    , ...,  306.    ,  113.    ,
                    6.4214],
                  [ -117.2  ,  32.77 ,  31.    , ...,  936.    ,  462.    ,
                    2.8621],
                  ...,
                  [ -116.4  ,  34.09 ,   9.    , ..., 2098.    ,  765.    ,
                    3.2723],
                  [ -118.01 ,  33.82 ,  31.    , ..., 1356.    ,  356.    ,
                    4.0625],
                  [ -122.45 ,  37.77 ,  52.    , ..., 1269.    ,  639.    ,
                    3.575 ]])
```

This array doesn't have any missing values now. We need to convert it into a dataframe again

```
In [104... type(X)
```

```
Out[104... numpy.ndarray
```

```
In [105... traindata_tr = pd.DataFrame(X, columns = traindata_num.columns)
```

```
In [106... traindata_tr
```

```
Out[106...
      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households  median_income
0      -121.89    37.29             38.0      1568.0          351.0         710.0         339.0          2.7042
1      -121.93    37.05             14.0       679.0          108.0         306.0         113.0          6.4214
2      -117.20    32.77             31.0     1952.0          471.0         936.0         462.0          2.8621
3      -119.61    36.31             25.0     1847.0          371.0        1460.0         353.0          1.8839
4      -118.59    34.23             17.0     6592.0         1525.0        4459.0        1463.0          3.0347
...      ...      ...             ...      ...      ...      ...      ...      ...
16507  -118.13    34.20             46.0     1271.0          236.0         573.0         210.0          4.9312
16508  -117.56    33.88             40.0     1196.0          294.0        1052.0         258.0          2.0682
16509  -116.40    34.09              9.0     4855.0          872.0        2098.0         765.0          3.2723
16510  -118.01    33.82             31.0     1960.0          380.0        1356.0         356.0          4.0625
16511  -122.45    37.77             52.0     3095.0          682.0        1269.0         639.0          3.5750
```

16512 rows × 8 columns

This dataframe doesn't have any missing values. Let's check

```
In [107... traindata_tr[traindata_tr.isna().any(axis = 1)].head()
```

```
Out[107...
      longitude  latitude  housing_median_age  total_rooms  total_bedrooms  population  households  median_income
```

This shows that there's no missing data now

```
In [108... traindata_tr.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16512 entries, 0 to 16511
Data columns (total 8 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude           16512 non-null  float64
1   latitude            16512 non-null  float64
2   housing_median_age  16512 non-null  float64
3   total_rooms         16512 non-null  float64
4   total_bedrooms      16512 non-null  float64
5   population          16512 non-null  float64
6   households          16512 non-null  float64
7   median_income       16512 non-null  float64
dtypes: float64(8)
memory usage: 1.0 MB
```

All the columns now have same number of data in the data fame.

```
In [109... traindata.head()
```

```
Out[109...
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN

```
In [110... #understanding the categorical data in the dataframe

traindata['ocean_proximity'].value_counts()
```

```
Out[110... <1H OCEAN    7276
INLAND        5263
NEAR OCEAN    2124
NEAR BAY      1847
ISLAND         2
Name: ocean_proximity, dtype: int64
```

```
In [111... traindata_cat = traindata['ocean_proximity']
```

```
In [112... # converting oceanproximity to nos.
traindata_cat_encoded,traindata_categories = traindata_cat.factorize()
```

```
In [113... traindata_cat_encoded[:10]
```

```
Out[113... array([0, 0, 1, 2, 0, 2, 0, 2, 0, 0], dtype=int64)
```

```
In [114... traindata_categories
```

```
Out[114... Index(['<1H OCEAN', 'NEAR OCEAN', 'INLAND', 'NEAR BAY', 'ISLAND'], dtype='object')
```

```
In [115... # Pandas factorize() example

df = pd.DataFrame({
    'A':['type1','type3','type3', 'type2', 'type0']
})
df['A'].factorize()
```

```
Out[115... (array([0, 1, 1, 2, 3], dtype=int64),
Index(['type1', 'type3', 'type2', 'type0'], dtype='object'))
```

```
In [116... traindata_cat.head(10)
```

```
Out[116... 17606    <1H OCEAN
18632    <1H OCEAN
14650    NEAR OCEAN
3230      INLAND
3555    <1H OCEAN
19480      INLAND
8879    <1H OCEAN
13685      INLAND
4937    <1H OCEAN
4861    <1H OCEAN
Name: ocean_proximity, dtype: object
```

```
In [117... traindata_cat_encoded[:20]
```

```
Out[117... array([0, 0, 1, 2, 0, 2, 0, 0, 2, 2, 0, 2, 2, 0, 3, 2, 2, 2],
      dtype=int64)
```

```
In [118... df1 =pd.DataFrame({'B':['okay', 'good', 'better', 'best']})
df1['B'].factorize()
```

```
Out[118... (array([0, 1, 2, 3], dtype=int64),
      Index(['okay', 'good', 'better', 'best'], dtype='object'))
```

```
In [119... from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder()
```

```
In [120... traindata_cat_1hot = encoder.fit_transform(traindata_cat_encoded.reshape(-1,1))
```

```
In [121... traindata_cat_1hot
```

```
Out[121... <16512x5 sparse matrix of type '<class 'numpy.float64'>'
      with 16512 stored elements in Compressed Sparse Row format>
```

```
In [122... type(traindata_cat_1hot)
```

```
Out[122... scipy.sparse.csr.csr_matrix
```

```
In [123... traindata_cat_1hot.toarray()
```

```
Out[123... array([[1., 0., 0., 0., 0.],
       [1., 0., 0., 0., 0.],
       [0., 1., 0., 0., 0.],
       ...,
       [0., 0., 1., 0., 0.],
       [1., 0., 0., 0., 0.],
       [0., 0., 0., 1., 0.]])
```

```
In [124... traindata_cat_encoded.reshape(-1,1)
```

```
Out[124... array([[0],
       [0],
       [1],
       ...,
       [2],
       [0],
       [3]], dtype=int64)
```

```
In [125... # Just run this cell, or copy it to your code, do not try to understand it (yet).
# Definition of the CategoricalEncoder class, copied from PR #9151.
```

```
from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.utils import check_array
from sklearn.preprocessing import LabelEncoder
from scipy import sparse
```

```
class CategoricalEncoder(BaseEstimator, TransformerMixin):
    """Encode categorical features as a numeric array.
    The input to this transformer should be a matrix of integers or strings,
    denoting the values taken on by categorical (discrete) features.
    The features can be encoded using a one-hot aka one-of-K scheme
    (`encoding='onehot'`, the default) or converted to ordinal integers
    (`encoding='ordinal'`).
    This encoding is needed for feeding categorical data to many scikit-learn
    estimators, notably linear models and SVMs with the standard kernels.
    Read more in the :ref:`User Guide <preprocessing_categorical_features>`.
    Parameters
    -----
    encoding : str, 'onehot', 'onehot-dense' or 'ordinal'
        The type of encoding to use (default is 'onehot'):
        - 'onehot': encode the features using a one-hot aka one-of-K scheme
```

```

        (or also called 'dummy' encoding). This creates a binary column for
        each category and returns a sparse matrix.
    - 'onehot-dense': the same as 'onehot' but returns a dense array
      instead of a sparse matrix.
    - 'ordinal': encode the features as ordinal integers. This results in
      a single column of integers (0 to n_categories - 1) per feature.
categories : 'auto' or a list of lists/arrays of values.
    Categories (unique values) per feature:
    - 'auto' : Determine categories automatically from the training data.
    - list : ``categories[i]`` holds the categories expected in the ith
      column. The passed categories are sorted before encoding the data
      (used categories can be found in the ``categories_`` attribute).
dtype : number type, default np.float64
    Desired dtype of output.
handle_unknown : 'error' (default) or 'ignore'
    Whether to raise an error or ignore if a unknown categorical feature is
    present during transform (default is to raise). When this is parameter
    is set to 'ignore' and an unknown category is encountered during
    transform, the resulting one-hot encoded columns for this feature
    will be all zeros.
    Ignoring unknown categories is not supported for
    ``encoding='ordinal'``.
Attributes
-----
categories_ : list of arrays
    The categories of each feature determined during fitting. When
    categories were specified manually, this holds the sorted categories
    (in order corresponding with output of ``transform``).
Examples
-----
Given a dataset with three features and two samples, we let the encoder
find the maximum value per feature and transform the data to a binary
one-hot encoding.
>>> from sklearn.preprocessing import CategoricalEncoder
>>> enc = CategoricalEncoder(handle_unknown='ignore')
>>> enc.fit([[0, 0, 3], [1, 1, 0], [0, 2, 1], [1, 0, 2]])
... # doctest: +ELLIPSIS
CategoricalEncoder(categories='auto', dtype=<... 'numpy.float64'>,
                    encoding='onehot', handle_unknown='ignore')
>>> enc.transform([[0, 1, 1], [1, 0, 4]]).toarray()
array([[ 1.,  0.,  0.,  1.,  0.,  0.,  1.,  0.,  0.],
       [ 0.,  1.,  1.,  0.,  0.,  0.,  0.,  0.,  0.]])
See also
-----
sklearn.preprocessing.OneHotEncoder : performs a one-hot encoding of
integer ordinal features. The ``OneHotEncoder`` assumes ``that input
features take on values in the range ``[0, max(feature)]`` instead of
using the unique values.
sklearn.feature_extraction.DictVectorizer : performs a one-hot encoding of
dictionary items (also handles string-valued features).
sklearn.feature_extraction.FeatureHasher : performs an approximate one-hot
encoding of dictionary items or strings.
"""

def __init__(self, encoding='onehot', categories='auto', dtype=np.float64,
             handle_unknown='error'):
    self.encoding = encoding
    self.categories = categories
    self.dtype = dtype
    self.handle_unknown = handle_unknown

def fit(self, X, y=None):
    """Fit the CategoricalEncoder to X.
    Parameters
    -----
    X : array-like, shape [n_samples, n_feature]
        The data to determine the categories of each feature.
    Returns
    -----
    self
    """

    if self.encoding not in ['onehot', 'onehot-dense', 'ordinal']:
        template = ("encoding should be either 'onehot', 'onehot-dense' "
                    "or 'ordinal', got %s")
        raise ValueError(template % self.handle_unknown)

    if self.handle_unknown not in ['error', 'ignore']:
        template = ("handle_unknown should be either 'error' or "
                    "'ignore', got %s")
        raise ValueError(template % self.handle_unknown)

    if self.encoding == 'ordinal' and self.handle_unknown == 'ignore':
        raise ValueError("handle_unknown='ignore' is not supported for "
                         "encoding='ordinal'")

    X = check_array(X, dtype=np.object, accept_sparse='csc', copy=True)
    n_samples, n_features = X.shape

```



```

self._label_encoders_ = [LabelEncoder() for i in range(n_features)]

for i in range(n_features):
    le = self._label_encoders_[i]
    Xi = X[:, i]
    if self.categories == 'auto':
        le.fit(Xi)
    else:
        valid_mask = np.in1d(Xi, self.categories[i])
        if not np.all(valid_mask):
            if self.handle_unknown == 'error':
                diff = np.unique(Xi[~valid_mask])
                msg = ("Found unknown categories {0} in column {1}"
                       " during fit".format(diff, i))
                raise ValueError(msg)
            le.classes_ = np.array(np.sort(self.categories[i]))

self.categories_ = [le.classes_ for le in self._label_encoders_]

return self

def transform(self, X):
    """Transform X using one-hot encoding.
    Parameters
    -----
    X : array-like, shape [n_samples, n_features]
        The data to encode.
    Returns
    -----
    X_out : sparse matrix or a 2-d array
        Transformed input.
    """
    X = check_array(X, accept_sparse='csc', dtype=np.object, copy=True)
    n_samples, n_features = X.shape
    X_int = np.zeros_like(X, dtype=np.int)
    X_mask = np.ones_like(X, dtype=np.bool)

    for i in range(n_features):
        valid_mask = np.in1d(X[:, i], self.categories_[i])

        if not np.all(valid_mask):
            if self.handle_unknown == 'error':
                diff = np.unique(X[~valid_mask, i])
                msg = ("Found unknown categories {0} in column {1}"
                       " during transform".format(diff, i))
                raise ValueError(msg)
            else:
                # Set the problematic rows to an acceptable value and
                # continue `The rows are marked `X_mask` and will be
                # removed later.
                X_mask[:, i] = valid_mask
                X[:, i][~valid_mask] = self.categories_[i][0]
            X_int[:, i] = self._label_encoders_[i].transform(X[:, i])

    if self.encoding == 'ordinal':
        return X_int.astype(self.dtype, copy=False)

    mask = X_mask.ravel()
    n_values = [cats.shape[0] for cats in self.categories_]
    n_values = np.array([0] + n_values)
    indices = np.cumsum(n_values)

    column_indices = (X_int + indices[:-1]).ravel()[mask]
    row_indices = np.repeat(np.arange(n_samples, dtype=np.int32),
                             n_features)[mask]
    data = np.ones(n_samples * n_features)[mask]

    out = sparse.csc_matrix((data, (row_indices, column_indices)),
                             shape=(n_samples, indices[-1]),
                             dtype=self.dtype).tocsr()

    if self.encoding == 'onehot-dense':
        return out.toarray()
    else:
        return out

```

In [126..

```

# Creating custom Transformer
from sklearn.base import BaseEstimator, TransformerMixin

# column index
rooms_ix, bedrooms_ix, population_ix, household_ix = 3, 4, 5, 6

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
    def __init__(self, add_bedrooms_per_room = True): # no *args or **kargs
        self.add_bedrooms_per_room = add_bedrooms_per_room
    def fit(self, X, y=None):
        return self # nothing else to do
    def transform(self, X, y=None):

```

```

rooms_per_household = X[:, rooms_ix] / X[:, household_ix]
population_per_household = X[:, population_ix] / X[:, household_ix]
if self.add_bedrooms_per_room:
    bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
    return np.c_[X, rooms_per_household, population_per_household,
                  bedrooms_per_room]
else:
    return np.c_[X, rooms_per_household, population_per_household]

attr_adder = CombinedAttributesAdder(add_bedrooms_per_room=False)
traindata_extra_attribs = attr_adder.transform(traindata.values)
traindata_extra_attribs = pd.DataFrame(traindata_extra_attribs, columns=list(traindata.columns)+["rooms_per_household"])
traindata_extra_attribs.head()

```

```

Out[126...

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	rooms_per_household
0	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN	4.920353982300885
1	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN	5.999999999999999
2	-117.2	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN	4.222943722943723
3	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND	5.23229461756374
4	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN	4.162747847573554

```

In [127...
max(traindata['total_rooms']), min(traindata['total_rooms']))

```

```

Out[127...
(39320.0, 6.0)

```

```


In [128...
plt.plot(traindata['total_rooms'], traindata['median_income'])

```

```

Out[128...
[<matplotlib.lines.Line2D at 0x1a68ed99fd0>]

```



```


In [129...
plt.scatter(traindata['total_rooms'], traindata['median_income'])

```

```

Out[129...
<matplotlib.collections.PathCollection at 0x1a68b049f10>

```



```

In [130...
# Feature Scaling - Min-max Scaling - Example
# Creating DataFrame first

s1 = pd.Series([1, 2, 3, 4, 5, 6], index=(range(6)))

```

```
s2 = pd.Series([10, 9, 8, 7, 6, 5], index=(range(6)))
df = pd.DataFrame(s1, columns=['s1'])
df['s2'] = s2
df
```

```
Out[130...]
   s1  s2
0    1  10
1    2   9
2    3   8
3    4   7
4    5   6
5    6   5
```

```
In [131...]
# Use Scikit-Learn minmax_scaling

#from mlxtend.preprocessing import minmax_scaling
#minmax_scaling(df, columns=['s1', 's2'])
```

```
In [132...]
#standardisation of the data
from sklearn.preprocessing import StandardScaler
stdscaler = StandardScaler()
df_tr = stdscaler.fit_transform(df)
df_tr
```

```
Out[132...]
array([[ -1.46385011,   1.46385011],
       [-0.87831007,   0.87831007],
       [-0.29277002,   0.29277002],
       [ 0.29277002,  -0.29277002],
       [ 0.87831007,  -0.87831007],
       [ 1.46385011,  -1.46385011]])
```

```
In [133...]
pd.DataFrame(df_tr, columns=['s1', 's2'])
```

```
Out[133...]
   s1      s2
0 -1.46385  1.46385
1 -0.87831  0.87831
2 -0.29277  0.29277
3  0.29277 -0.29277
4  0.87831 -0.87831
5  1.46385 -1.46385
```

```
In [134...]
#performing a series of transformation using pipeline
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
#note: we have already created a method by the name CombinedAttributesAdder() above and defined fit and transform
num_pipeline = Pipeline([('imputer', SimpleImputer(strategy='median')), ('attrs_adder', CombinedAttributesAdder())])
traindata_num_tr = num_pipeline.fit_transform(traindata_num)
```

```
In [135...]
traindata_num_tr
# note: scikit learn does not handle dataframe
```

```
Out[135...]
array([[ -1.15604281,   0.77194962,   0.74333089, ..., -0.31205452,
        -0.08649871,   0.15531753],
       [-1.17602483,   0.6596948 , -1.1653172 , ...,   0.21768338,
        -0.03353391,  -0.83628902],
       [ 1.18684903, -1.34218285,   0.18664186, ..., -0.46531516,
        -0.09240499,   0.4222004 ],
       ...,
       [ 1.58648943, -0.72478134, -1.56295222, ...,   0.3469342 ,
        -0.03055414, -0.52177644],
       [ 0.78221312, -0.85106801,   0.18664186, ...,   0.02499488,
        0.06150916, -0.30340741],
       [-1.43579109,   0.99645926,   1.85670895, ..., -0.22852947,
        -0.09586294,   0.10180567]])
```

```
In [136... traindata_num
```

```
Out[136... longitude latitude housing_median_age total_rooms total_bedrooms population households median_income
```

17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347
...
6563	-118.13	34.20	46.0	1271.0	236.0	573.0	210.0	4.9312
12053	-117.56	33.88	40.0	1196.0	294.0	1052.0	258.0	2.0682
13908	-116.40	34.09	9.0	4855.0	872.0	2098.0	765.0	3.2723
11159	-118.01	33.82	31.0	1960.0	380.0	1356.0	356.0	4.0625
15775	-122.45	37.77	52.0	3095.0	682.0	1269.0	639.0	3.5750

16512 rows × 8 columns

```
In [137... type(traindata_num_tr)
```

```
Out[137... numpy.ndarray
```

```
In [138... num_attribs = list(traindata_num)
num_attribs
```

```
Out[138... ['longitude',
'latitude',
'housing_median_age',
'total_rooms',
'total_bedrooms',
'population',
'households',
'median_income']
```

```
In [139... from sklearn.base import BaseEstimator, TransformerMixin

# Create a class to select numerical or categorical columns
# since Scikit-Learn doesn't handle DataFrames yet
class DataFrameSelector(BaseEstimator, TransformerMixin):
    def __init__(self, attribute_names):
        self.attribute_names = attribute_names
    def fit(self, X, y=None):
        return self
    def transform(self, X):
        return X[self.attribute_names].values
```

```
In [140... num_attribs = list(traindata_num)
num_pipeline = Pipeline([('selector', DataFrameSelector(num_attribs)),('imputer', SimpleImputer(strategy='median'))])
traindata_num_tr = num_pipeline.fit_transform(traindata_num)

#we've also created a pipeline for categorical attributes
cat_attribs = ['ocean_proximity']
cat_pipeline = Pipeline([('selector', DataFrameSelector(cat_attribs)),
                          ('cat_encoder', CategoricalEncoder(encoding='onehot-dense'))])
```

```
In [141... Y = num_pipeline.fit_transform(traindata)
# This is the command to fit and transform the numerical attributes. since scikit learn does not yeild dataframe
#retured
```

```
In [142... type(Y)
```

```
Out[142... numpy.ndarray
```

In [143...

Y

Out[143...

```
array([[ -1.15604281,  0.77194962,  0.74333089, ..., -0.31205452,
        -0.08649871,  0.15531753],
       [ -1.17602483,  0.6596948 , -1.1653172 , ...,  0.21768338,
        -0.03353391, -0.83628902],
       [  1.18684903, -1.34218285,  0.18664186, ..., -0.46531516,
        -0.09240499,  0.4222004 ],
       ...,
       [  1.58648943, -0.72478134, -1.56295222, ...,  0.3469342 ,
        -0.03055414, -0.52177644],
       [  0.78221312, -0.85106801,  0.18664186, ...,  0.02499488,
        0.06150916, -0.30340741],
       [-1.43579109,  0.99645926,  1.85670895, ..., -0.22852947,
        -0.09586294,  0.10180567]])
```

In [144...

```
#converting the array into data frame
#note: we have used the columns of the dataframe that has all the extra attributed added to it. we cant use the c
# and the new extra added attributes dataframe contains 11 columns
traindata_scaled_tr = pd.DataFrame(Y, columns =traindata_extra_attribs.columns)
traindata_scaled_tr
```

Out[144...

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	room
0	-1.156043	0.771950	0.743331	-0.493234	-0.445438	-0.636211	-0.420698	-0.614937	-0.312055	
1	-1.176025	0.659695	-1.165317	-0.908967	-1.036928	-0.998331	-1.022227	1.336459	0.217683	
2	1.186849	-1.342183	0.186642	-0.313660	-0.153345	-0.433639	-0.093318	-0.532046	-0.465315	
3	-0.017068	0.313576	-0.290520	-0.362762	-0.396756	0.036041	-0.383436	-1.045566	-0.079661	
4	0.492474	-0.659299	-0.926736	1.856193	2.412211	2.724154	2.570975	-0.441437	-0.357834	
...
16507	0.722267	-0.673331	1.379547	-0.632123	-0.725361	-0.759010	-0.764049	0.554158	0.234352	
16508	1.007011	-0.823004	0.902385	-0.667196	-0.584183	-0.329664	-0.636291	-0.948815	-0.308114	
16509	1.586489	-0.724781	-1.562952	1.043901	0.822735	0.607904	0.713156	-0.316705	0.346934	
16510	0.782213	-0.851068	0.186642	-0.309919	-0.374849	-0.057178	-0.375451	0.098121	0.024995	
16511	-1.435791	0.996459	1.856709	0.220853	0.360253	-0.135159	0.377791	-0.157799	-0.228529	

16512 rows × 11 columns

In [145...

```
#we are not going to run the pipeline for the categetical data seperately. Instead we will use the FeatureUnion (
#concatenation of the pipelines
from sklearn.pipeline import FeatureUnion

full_pipeline = FeatureUnion(transformer_list=[
    ("num_pipeline", num_pipeline),
    ("cat_pipeline", cat_pipeline),
])
```

In [146...

type(full_pipeline)

Out[146...

sklearn.pipeline.FeatureUnion

In [147...

```
traindata_prepared_arr = full_pipeline.fit_transform(traindata)
#this is the array with numerical as well as categorical attributes
traindata_prepared_arr[0]
traindata_prepared_arr[1]
```

```
<ipython-input-125-341a9c0cd727>:110: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
X = check_array(X, dtype=np.object, accept_sparse='csc', copy=True)
<ipython-input-125-341a9c0cd727>:145: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
```

```

X = check_array(X, accept_sparse='csc', dtype=np.object, copy=True)
<ipython-input-125-341a9c0cd727>:147: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To
silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing
`np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your
current use, check the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#depr
ecations
X_int = np.zeros_like(X, dtype=np.int)
<ipython-input-125-341a9c0cd727>:148: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`.
To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specif
ically wanted the numpy scalar type, use `np.bool_` here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#depr
ecations
X_mask = np.ones_like(X, dtype=np.bool)

```

```

Out[147]: array([-1.17602483,  0.6596948 , -1.1653172 , -0.90896655, -1.0369278 ,
                -0.99833135, -1.02222705,  1.33645936,  0.21768338, -0.03353391,
                -0.83628902,  1.          ,  0.          ,  0.          ,  0.          ,
                0.          ])

```

```

In [148]: traindata_prepared_arr[1]

```

```

Out[148]: array([-1.17602483,  0.6596948 , -1.1653172 , -0.90896655, -1.0369278 ,
                -0.99833135, -1.02222705,  1.33645936,  0.21768338, -0.03353391,
                -0.83628902,  1.          ,  0.          ,  0.          ,  0.          ,
                0.          ])

```

```

In [149]: traindata_prepared_arr[52]

```

```

Out[149]: array([-1.18601584,  0.75791776,  0.42522288, -0.53345104, -0.67911311,
                -0.7491498 , -0.66823011,  0.85910944,  0.22363394, -0.06236395,
                -0.63722671,  1.          ,  0.          ,  0.          ,  0.          ,
                0.          ])

```

```

In [150]: traindata_prepared_arr

```

```

Out[150]: array([[ -1.15604281,  0.77194962,  0.74333089, ...,  0.          ,
                  0.          ,  0.          ],
                [ -1.17602483,  0.6596948 , -1.1653172 , ...,  0.          ,
                  0.          ,  0.          ],
                [  1.18684903, -1.34218285,  0.18664186, ...,  0.          ,
                  0.          ,  1.          ],
                ...,
                [  1.58648943, -0.72478134, -1.56295222, ...,  0.          ,
                  0.          ,  0.          ],
                [  0.78221312, -0.85106801,  0.18664186, ...,  0.          ,
                  0.          ,  0.          ],
                [-1.43579109,  0.99645926,  1.85670895, ...,  0.          ,
                  1.          ,  0.          ]])

```

```

In [151]: traindata_labels.head()
          type(traindata_labels)

```

```

Out[151]: pandas.core.series.Series

```

```

In [152]: traindata_scaled_tr

```

```

Out[152]:

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity	room
0	-1.156043	0.771950	0.743331	-0.493234	-0.445438	-0.636211	-0.420698	-0.614937	-0.312055	
1	-1.176025	0.659695	-1.165317	-0.908967	-1.036928	-0.998331	-1.022227	1.336459	0.217683	
2	1.186849	-1.342183	0.186642	-0.313660	-0.153345	-0.433639	-0.093318	-0.532046	-0.465315	
3	-0.017068	0.313576	-0.290520	-0.362762	-0.396756	0.036041	-0.383436	-1.045566	-0.079661	
4	0.492474	-0.659299	-0.926736	1.856193	2.412211	2.724154	2.570975	-0.441437	-0.357834	

...
16507	0.722267	-0.673331	1.379547	-0.632123	-0.725361	-0.759010	-0.764049	0.554158	0.234352
16508	1.007011	-0.823004	0.902385	-0.667196	-0.584183	-0.329664	-0.636291	-0.948815	-0.308114
16509	1.586489	-0.724781	-1.562952	1.043901	0.822735	0.607904	0.713156	-0.316705	0.346934
16510	0.782213	-0.851068	0.186642	-0.309919	-0.374849	-0.057178	-0.375451	0.098121	0.024995
16511	-1.435791	0.996459	1.856709	0.220853	0.360253	-0.135159	0.377791	-0.157799	-0.228529

16512 rows × 11 columns

--	--	--	--	--	--	--	--	--	--	--

In [153...

```
traindata
```

Out[153...

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
17606	-121.89	37.29	38.0	1568.0	351.0	710.0	339.0	2.7042	<1H OCEAN
18632	-121.93	37.05	14.0	679.0	108.0	306.0	113.0	6.4214	<1H OCEAN
14650	-117.20	32.77	31.0	1952.0	471.0	936.0	462.0	2.8621	NEAR OCEAN
3230	-119.61	36.31	25.0	1847.0	371.0	1460.0	353.0	1.8839	INLAND
3555	-118.59	34.23	17.0	6592.0	1525.0	4459.0	1463.0	3.0347	<1H OCEAN
...
6563	-118.13	34.20	46.0	1271.0	236.0	573.0	210.0	4.9312	INLAND
12053	-117.56	33.88	40.0	1196.0	294.0	1052.0	258.0	2.0682	INLAND
13908	-116.40	34.09	9.0	4855.0	872.0	2098.0	765.0	3.2723	INLAND
11159	-118.01	33.82	31.0	1960.0	380.0	1356.0	356.0	4.0625	<1H OCEAN
15775	-122.45	37.77	52.0	3095.0	682.0	1269.0	639.0	3.5750	NEAR BAY

16512 rows × 9 columns

In [154...

```
#training the model-Linear Regression
from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression() #initializing the object
h = lin_reg.fit(traindata_prepared_arr, traindata_labels)
```

In [155...

```
type(h)
```

Out[155...

```
sklearn.linear_model._base.LinearRegression
```

In [156...

```
h
```

Out[156...

```
LinearRegression()
```

In [157...

```
#lets try the full pipeline on new instances
some_data = traindata.iloc[:5]
some_labels = traindata_labels[:5]
some_data_prepared_arr = full_pipeline.transform(some_data)
```

```
<ipython-input-125-341a9c0cd727>:145: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
X = check_array(X, accept_sparse='csc', dtype=np.object, copy=True)
<ipython-input-125-341a9c0cd727>:147: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
X_int = np.zeros_like(X, dtype=np.int)
<ipython-input-125-341a9c0cd727>:148: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
```

```
X_mask = np.ones_like(X, dtype=np.bool)
```

```
In [158... some_data_prepared_arr
```

```
Out[158... array([[ -1.15604281,  0.77194962,  0.74333089, -0.49323393, -0.44543821,
          -0.63621141, -0.42069842, -0.61493744, -0.31205452, -0.08649871,
           0.15531753,  1.          ,  0.          ,  0.          ,  0.          ,
           0.          ],
        [ -1.17602483,  0.6596948 , -1.1653172 , -0.90896655, -1.0369278 ,
          -0.99833135, -1.02222705,  1.33645936,  0.21768338, -0.03353391,
          -0.83628902,  1.          ,  0.          ,  0.          ,  0.          ,
           0.          ],
        [  1.18684903, -1.34218285,  0.18664186, -0.31365989, -0.15334458,
          -0.43363936, -0.0933178 , -0.5320456 , -0.46531516, -0.09240499,
           0.4222004 ,  0.          ,  0.          ,  0.          ,  0.          ,
           1.          ],
        [ -0.01706767,  0.31357576, -0.29052016, -0.36276217, -0.39675594,
           0.03604096, -0.38343559, -1.04556555, -0.07966124,  0.08973561,
          -0.19645314,  0.          ,  1.          ,  0.          ,  0.          ,
           0.          ],
        [  0.49247384, -0.65929936, -0.92673619,  1.85619316,  2.41221109,
           2.72415407,  2.57097492, -0.44143679, -0.35783383, -0.00419445,
           0.2699277 ,  1.          ,  0.          ,  0.          ,  0.          ,
           0.          ]])
```

Understanding the algorithm in Detail

We extracted five rows from the trained dataframe containing all the attributes except median_house valuesome

Then we extracted five rows from the median_house values

In the third step:

We sent the dataframe 'some_data' to the 'full_pipeline' which contains the object FeatureUnion that combined the categorical and numerical pipeline together. The dataframe 'some_data' goes through all the cleaning necessary through these pipelines and then is transformed into an array containing all its attributes.

```
In [159... #now we do the prediction
print('predictions', h.predict(some_data_prepared_arr))
```

```
predictions [210644.60459286 317768.80697211 210956.43331178  59218.98886849
 189747.55849879]
```

```
In [160... print('labels', list(some_labels))
```

```
labels [286600.0, 340600.0, 196900.0, 46300.0, 254500.0]
```

```
In [161... traindata_predictions = h.predict(traindata_prepared_arr) # this predicts all the data
```

```
In [162... type(traindata_predictions)
```

```
Out[162... numpy.ndarray
```

```
In [163... traindata_predictions # this is an array
len(traindata_predictions)
```

```
Out[163... 16512
```

```
In [164... type(traindata_labels) # this is a series
len(traindata_labels)
```

```
Out[164... 16512
```



```
In [165... #calculated mean squared error
from sklearn.metrics import mean_squared_error
lin_mse = mean_squared_error(traindata_predictions, traindata_labels)
type(lin_mse)
```

```
Out[165... numpy.float64
```

```
In [166... lin_mse
```

```
Out[166... 4709829587.971121
```

```
In [167... lin_rmse = np.sqrt(lin_mse)
lin_rmse
```

```
Out[167... 68628.19819848923
```

The error is huge indeed. This means that the fitted has deviated from the actual value by about 68628 above or below. This is an example of underfitting.

```
In [168... #using decision trees on the data
#this is another model used for predictions
from sklearn.tree import DecisionTreeRegressor
tree_reg = DecisionTreeRegressor(random_state=42) # initilisation
tree_reg.fit(traindata_prepared_arr, traindata_labels) #this commands help the algorithm learn from the data we
```

```
Out[168... DecisionTreeRegressor(random_state=42)
```

```
In [169... #now we predict from the fitted data, it will generate an array of fitted values of housing_median_values
traindata_predictions_1 = tree_reg.predict(traindata_prepared_arr)
```

```
In [170... traindata_predictions_1[:5]
```

```
Out[170... array([286600., 340600., 196900., 46300., 254500.])
```

```
In [171... #caculate the root mean squared error..
tree_mse = mean_squared_error(traindata_labels, traindata_predictions_1)
tree_rmse = np.sqrt(tree_mse)
tree_rmse
```

```
Out[171... 0.0
```

The RMSE obtained is zero. This shows that the model seems to have memorized the whole training set. This is called overfitting. So we perform the Cross-validation over the training set without touching the training set. We shall create training folds out of the training set

```
In [172... from sklearn.model_selection import cross_val_score
tree_reg = DecisionTreeRegressor(random_state=42)
scores = cross_val_score(tree_reg, traindata_prepared_arr, traindata_labels, scoring= 'neg_mean_squared_error', cv=5)
```

```
In [173... scores
```

```
Out[173... array([-4.92724492e+09, -4.46961291e+09, -5.24647900e+09, -5.00679914e+09,
        -5.05746872e+09, -5.71311365e+09, -4.93686969e+09, -4.93838343e+09,
        -5.68016653e+09, -5.07394900e+09])
```

```
In [174... tree_rmse_scores = np.sqrt(-scores)
```

```
tree_rmse_scores
```

```
Out[174]: array([70194.33680785, 66855.16363941, 72432.58244769, 70758.73896782,
        71115.88230639, 75585.14172901, 70262.86139133, 70273.6325285 ,
        75366.87952553, 71231.65726027])
```

```
In [175]: #Look at the score of cross-validation of DecisionTreeRegressor
```

```
def display_scores(scores):
    print("Scores:", scores)
    print("Mean:", scores.mean())
    print("Standard deviation:", scores.std())
```

```
display_scores(tree_rmse_scores)
```

```
Scores: [70194.33680785 66855.16363941 72432.58244769 70758.73896782
        71115.88230639 75585.14172901 70262.86139133 70273.6325285
        75366.87952553 71231.65726027]
Mean: 71407.68766037929
Standard deviation: 2439.4345041191004
```

```
In [176]: #using cross validation on Linear Regression
```

```
lin_reg = LinearRegression()
lin_scores = cross_val_score(lin_reg, traindata_prepared_arr, traindata_labels, scoring= 'neg_mean_squared_error')
lin_rmse_scores = np.sqrt(-lin_scores)
```

```
Out[176]: array([66782.73843989, 66960.118071 , 70347.95244419, 74739.57052552,
        68031.13388938, 71193.84183426, 64969.63056405, 68281.61137997,
        71552.91566558, 67665.10082067])
```

```
In [177]: display_scores(lin_rmse_scores)
```

```
Scores: [66782.73843989 66960.118071 70347.95244419 74739.57052552
        68031.13388938 71193.84183426 64969.63056405 68281.61137997
        71552.91566558 67665.10082067]
Mean: 69052.46136345083
Standard deviation: 2731.674001798348
```

So we see that the mean value(69052.46136345083) for Linear regression is less than the mean value(71407.68766037929) of the Decision Trees. Which one to choose? Let's see further

```
In [178]: # Let's train one more model using Random Forests
```

```
from sklearn.ensemble import RandomForestRegressor

forest_reg = RandomForestRegressor(random_state=42)
forest_reg.fit(traindata_prepared_arr, traindata_labels)
```

```
Out[178]: RandomForestRegressor(random_state=42)
```

```
In [179]: #prediction and MSE calculation
```

```
traindata_predictions_2 = forest_reg.predict(traindata_prepared_arr) #creates an array of predicted values
forest_mse = mean_squared_error(traindata_predictions_2, traindata_labels) #gives the mean squared error
forest_rmse = np.sqrt(forest_mse)
```

```
Out[179]: 18603.515021376355
```

```
In [180]: #cross validation for Random Forest
```

```
forest_reg = RandomForestRegressor(random_state=42)
forest_scores = cross_val_score(forest_reg, traindata_prepared_arr, traindata_labels, scoring= 'neg_mean_squared_error')
forest_rmse_scores = np.sqrt(-forest_scores)
```

```
In [181]: forest_rmse_scores
```

```
Out[181] array([49519.80364233, 47461.9115823 , 50029.02762854, 52325.28068953,
        49308.39426421, 53446.37892622, 48634.8036574 , 47585.73832311,
        53490.10699751, 50021.5852922 ])
```

```
In [182] display_scores(forest_rmse_scores)
```

```
Scores: [49519.80364233 47461.9115823  50029.02762854 52325.28068953
 49308.39426421 53446.37892622 48634.8036574  47585.73832311
 53490.10699751 50021.5852922 ]
Mean: 50182.303100336096
Standard deviation: 2097.0810550985693
```

so we have applied three models till now.

Now that we have checked three models, it is time to fine tune the model. In the process of fine tuning, we need to combine hyperparameters and perform gridsearching to get the best model. See copy

```
In [183] from sklearn.model_selection import GridSearchCV
# try 12 (3x4) combinations of hyperparameters
param_grid = [{'n_estimators': [3, 10, 30], 'max_features': [2, 4, 6, 8]},
               # then try 6 (2x3) combinations with bootstrap set as False
               {'bootstrap': [False], 'n_estimators': [3, 10], 'max_features': [2, 3, 4]}]
```

```
In [184] forest_reg = RandomForestRegressor(random_state = 42)
grid_search = GridSearchCV(forest_reg, param_grid, cv=5, scoring= 'neg_mean_squared_error')
grid_search.fit(traindata_prepared_arr, traindata_labels)
```

```
Out[184] GridSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
        param_grid=[{'max_features': [2, 4, 6, 8],
                      'n_estimators': [3, 10, 30]},
                     {'bootstrap': [False], 'max_features': [2, 3, 4],
                      'n_estimators': [3, 10]}],
        scoring='neg_mean_squared_error')
```

```
In [185] #best hyperparameter
grid_search.best_params_
```

```
Out[185] {'max_features': 8, 'n_estimators': 30}
```

```
In [186] #best estimator
grid_search.best_estimator_
```

```
Out[186] RandomForestRegressor(max_features=8, n_estimators=30, random_state=42)
```

```
In [187] #looking at score of each hyperparamter combination tested during the grid search
cvres = grid_search.cv_results_
for mean_score, params in zip(cvres['mean_test_score'], cvres['params']):
    print(np.sqrt(-mean_score), params)
```

```
63669.11631261028 {'max_features': 2, 'n_estimators': 3}
55627.099719926795 {'max_features': 2, 'n_estimators': 10}
53384.57275149205 {'max_features': 2, 'n_estimators': 30}
60965.950449450494 {'max_features': 4, 'n_estimators': 3}
52741.04704299915 {'max_features': 4, 'n_estimators': 10}
50377.40461678399 {'max_features': 4, 'n_estimators': 30}
58663.93866579625 {'max_features': 6, 'n_estimators': 3}
52006.19873526564 {'max_features': 6, 'n_estimators': 10}
50146.51167415009 {'max_features': 6, 'n_estimators': 30}
57869.25276169646 {'max_features': 8, 'n_estimators': 3}
51711.127883959234 {'max_features': 8, 'n_estimators': 10}
49682.273345071546 {'max_features': 8, 'n_estimators': 30}
62895.06951262424 {'bootstrap': False, 'max_features': 2, 'n_estimators': 3}
54658.176157539405 {'bootstrap': False, 'max_features': 2, 'n_estimators': 10}
59470.40652318466 {'bootstrap': False, 'max_features': 3, 'n_estimators': 3}
52724.9822587892 {'bootstrap': False, 'max_features': 3, 'n_estimators': 10}
```

```
57490.5691951261 {'bootstrap': False, 'max_features': 4, 'n_estimators': 3}
51009.495668875716 {'bootstrap': False, 'max_features': 4, 'n_estimators': 10}
```

Thus the lowest score we have is for (8,30)

49682.273345071546 {'max_features': 8, 'n_estimators': 30}

```
In [188... # RandomizedSearchCV

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_distributions = {
    'n_estimators': randint(low=1, high=200),
    'max_features': randint(low=1, high=8),
}

forest_reg = RandomForestRegressor(random_state=42)
rnd_search = RandomizedSearchCV(forest_reg, param_distributions=param_distributions,
                                n_iter=10, cv=5, scoring='neg_mean_squared_error', random_state=42)
rnd_search.fit(traindata_prepared_arr, traindata_labels)
```

```
Out[188... RandomizedSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
                    param_distributions={'max_features': <scipy.stats._distn_infrastructure.rv_frozen object at 0x
000001A68EF8EC40>,
                                         'n_estimators': <scipy.stats._distn_infrastructure.rv_frozen object at 0x
000001A68EFAD700>},
                    random_state=42, scoring='neg_mean_squared_error')
```

```
In [189... rnd_search.best_estimator_
```

```
Out[189... RandomForestRegressor(max_features=7, n_estimators=180, random_state=42)
```

```
In [190... rnd_search.best_params_
```

```
Out[190... {'max_features': 7, 'n_estimators': 180}
```

```
In [191... cvres1 = rnd_search.cv_results_
for mean_score, params in zip(cvres1['mean_test_score'], cvres1['params']):
    print(np.sqrt(-mean_score), params)

49150.70756927707 {'max_features': 7, 'n_estimators': 180}
51389.889203389284 {'max_features': 5, 'n_estimators': 15}
50796.155224308866 {'max_features': 3, 'n_estimators': 72}
50835.13360315349 {'max_features': 5, 'n_estimators': 21}
49280.9449827171 {'max_features': 7, 'n_estimators': 122}
50774.90662363929 {'max_features': 3, 'n_estimators': 75}
50682.78888164288 {'max_features': 3, 'n_estimators': 88}
49608.99608105296 {'max_features': 5, 'n_estimators': 100}
50473.61930350219 {'max_features': 3, 'n_estimators': 150}
64429.84143294435 {'max_features': 5, 'n_estimators': 2}
```

Thus the lowest score we have is for (7,180)

49150.70756927707 {'max_features': 7, 'n_estimators': 180}

```
In [194... #see the importance score of each attribute using GridsearchCV
feature_importances = grid_search.best_estimator_.feature_importances_
extra_attribs = ["rooms_per_hhold", "pop_per_hhold", "bedrooms_per_room"]
cat_encoder = cat_pipeline.named_steps["cat_encoder"]
cat_one_hot_attribs = list(cat_encoder.categories_[0])
attributes = num_attribs + extra_attribs + cat_one_hot_attribs
sorted(zip(feature_importances, attributes), reverse=True)
```

```
Out[194... [(0.36615898061813423, 'median_income'),
```

```
(0.16478099356159054, 'INLAND'),
(0.10879295677551575, 'pop_per_hhold'),
(0.07334423551601243, 'longitude'),
(0.06290907048262032, 'latitude'),
(0.056419179181954014, 'rooms_per_hhold'),
(0.053351077347675815, 'bedrooms_per_room'),
(0.04114379847872964, 'housing_median_age'),
(0.014874280890402769, 'population'),
(0.014672685420543239, 'total_rooms'),
(0.014257599323407808, 'households'),
(0.014106483453584104, 'total_bedrooms'),
(0.010311488326303788, '<1H OCEAN'),
(0.0028564746373201584, 'NEAR OCEAN'),
(0.0019604155994780706, 'NEAR BAY'),
(6.0280386727366e-05, 'ISLAND')]
```

In [195...

```
# Evaluate model on the Test Set
```

```
final_model = grid_search.best_estimator_
X_test = strat_test_set.drop("median_house_value", axis=1)
y_test = strat_test_set["median_house_value"].copy()
X_test_prepared = full_pipeline.transform(X_test)
```

<ipython-input-125-341a9c0cd727>:145: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X = check_array(X, accept_sparse='csc', dtype=np.object, copy=True)
```

<ipython-input-125-341a9c0cd727>:147: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X_int = np.zeros_like(X, dtype=np.int)
```

<ipython-input-125-341a9c0cd727>:148: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X_mask = np.ones_like(X, dtype=np.bool)
```

In [196...

```
#command for data cleaning through the pipeline that was created earlier
X_test_prepared = full_pipeline.transform(X_test)
```

<ipython-input-125-341a9c0cd727>:145: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X = check_array(X, accept_sparse='csc', dtype=np.object, copy=True)
```

<ipython-input-125-341a9c0cd727>:147: DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To silence this warning, use `int` by itself. Doing this will not modify any behavior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64` or `np.int32` to specify the precision. If you wish to review your current use, check the release note link for additional information. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X_int = np.zeros_like(X, dtype=np.int)
```

<ipython-input-125-341a9c0cd727>:148: DeprecationWarning: `np.bool` is a deprecated alias for the builtin `bool`. To silence this warning, use `bool` by itself. Doing this will not modify any behavior and is safe. If you specifically wanted the numpy scalar type, use `np.bool_` here. Deprecated in NumPy 1.20; for more details and guidance: <https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
X_mask = np.ones_like(X, dtype=np.bool)
```

In [197...

```
#prediction using the RandomForestRegressor mode. will create and array of predicted values for median_house_val
final_predictions = final_model.predict(X_test_prepared)
```

In [198...

```
#calculating the RMSE
final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)
final_rmse
```

Out[198... 47730.22690385927

```
In [199] final_model
```

```
Out[199] RandomForestRegressor(max_features=8, n_estimators=30, random_state=42)
```

```
In [200] X_test
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	ocean_proximity
5241	-118.39	34.12	29.0	6447.0	1012.0	2184.0	960.0	8.2816	<1H OCEAN
10970	-117.86	33.77	39.0	4159.0	655.0	1669.0	651.0	4.6111	<1H OCEAN
20351	-119.05	34.21	27.0	4357.0	926.0	2110.0	876.0	3.0119	<1H OCEAN
6568	-118.15	34.20	52.0	1786.0	306.0	1018.0	322.0	4.1518	INLAND
13285	-117.68	34.07	32.0	1775.0	314.0	1067.0	302.0	4.0375	INLAND
...
20519	-121.53	38.58	33.0	4988.0	1169.0	2414.0	1075.0	1.9728	INLAND
17430	-120.44	34.65	30.0	2265.0	512.0	1402.0	471.0	1.9750	NEAR OCEAN
4019	-118.49	34.18	31.0	3073.0	674.0	1486.0	684.0	4.8984	<1H OCEAN
12107	-117.32	33.99	27.0	5464.0	850.0	2400.0	836.0	4.7110	INLAND
2398	-118.91	36.79	19.0	1616.0	324.0	187.0	80.0	3.7857	INLAND

4128 rows × 9 columns

```
In [201] y_test
```

```
Out[201] 5241    500001.0
10970    240300.0
20351    218200.0
6568     182100.0
13285    121300.0
...
20519     76400.0
17430    134000.0
4019     311700.0
12107    133500.0
2398     78600.0
Name: median_house_value, Length: 4128, dtype: float64
```

```
In [202] X_test_prepared[0:5]
```

```
Out[202] array([[ 0.59238393, -0.71074948,  0.02758786,  1.78838525,  1.16351084,
        0.68498857,  1.23217448,  2.31299771,  0.48830927, -0.07090847,
       -0.86820063,  1.          ,  0.          ,  0.          ,  0.          ,
        0.          ],
       [ 0.8571457 , -0.87445443,  0.8228579 ,  0.71842323,  0.29453231,
        0.22337528,  0.40973048,  0.38611673,  0.36310326, -0.04598303,
       -0.86028018,  1.          ,  0.          ,  0.          ,  0.          ,
        0.          ],
       [ 0.26268061, -0.66865392, -0.13146615,  0.8110161 ,  0.95417708,
        0.61865967,  1.00859747, -0.45340597, -0.17866074, -0.05936925,
       -0.01792937,  1.          ,  0.          ,  0.          ,  0.          ,
        0.          ],
       [ 0.71227605, -0.67333121,  1.85670895, -0.39128825, -0.55497332,
       -0.36013977, -0.46594615,  0.14500069,  0.04068081,  0.00561556,
       -0.64843204,  0.          ,  1.          ,  0.          ,  0.          ,
        0.          ],
       [ 0.94706479, -0.7341359 ,  0.26616887, -0.3964323 , -0.53550041,
       -0.31621928, -0.51917877,  0.08499728,  0.167383 ,  0.03769486,
       -0.56320765,  0.          ,  1.          ,  0.          ,  0.          ,
        0.          ]])
```

```
In [203] final_predictions[0:10]
```

```
Out[203...] array([495467.5, 262676.7, 235380., 211883.33333333,
      135516.66666667, 147776.66666667, 63540., 439026.9,
      106323.33333333, 100293.33333333])
```

```
In [204...] y_test.head(10)
```

```
Out[204...] 5241      500001.0
10970      240300.0
20351      218200.0
6568       182100.0
13285      121300.0
20552      120600.0
19989       72300.0
17049      500001.0
13692      98900.0
13916       82600.0
Name: median_house_value, dtype: float64
```

```
In [ ]:
```